

Software vulnerabilities: memory corruption

Classifying vulnerabilities

- Memory corruption (safety)
 - Integer Overflow
 - Buffer Overflow
 - Stack Overflow
 - Use after free/ Double free
 - Null pointer dereference
- Exploiting memory corruption to control program
 - Format string attacks
 - Stack buffer overflow (or stack-based buffer overflow)
 - Heap buffer overflow (or heap-based buffer overflow)
 - Overflowing on non “control-data” memory
- Resulting attacks
 - Code injection into memory (e.g., shellcodes)
 - Code reuse techniques (Return-to-libc, Return-Oriented Programming, Jump-Oriented Programming)

Integer overflows

- Unsigned ints
 - 32 bits: range from 0 to $2^{32}-1$ (4,294,967,295)
 - Overflow: $4,294,967,295 + 1 = 0$
- Signed ints
 - 32 bits: range from $-(2^{31})$ to $2^{31} - 1$ (2,147,483,647)
 - Overflow: $2,147,483,647 + 2 = -2,147,483,648$
- Some languages (Java, Ada) throw exceptions, many don't
- Attacker can supply large values used in :
 - Computation of the size of a buffer allocation (malloc)
 - Array access (in particular the bound checks)
- Solution: always check overflow before critical operation!

Buffer Overflows

- A buffer overflow occurs any time a program attempts to store data beyond the boundaries of a buffer, overwriting the adjacent memory locations
- Originates from mistakes done while writing code
 - unfamiliarity with language
 - Boundary or arithmetic errors
- Mostly C / C++ programs
- Addressed by languages with automatic memory management: dynamic bounds checks (e.g., Java) or automatic resizing of buffers (e.g., Perl)
 - Still native libraries that are written in C (e.g., JNI)

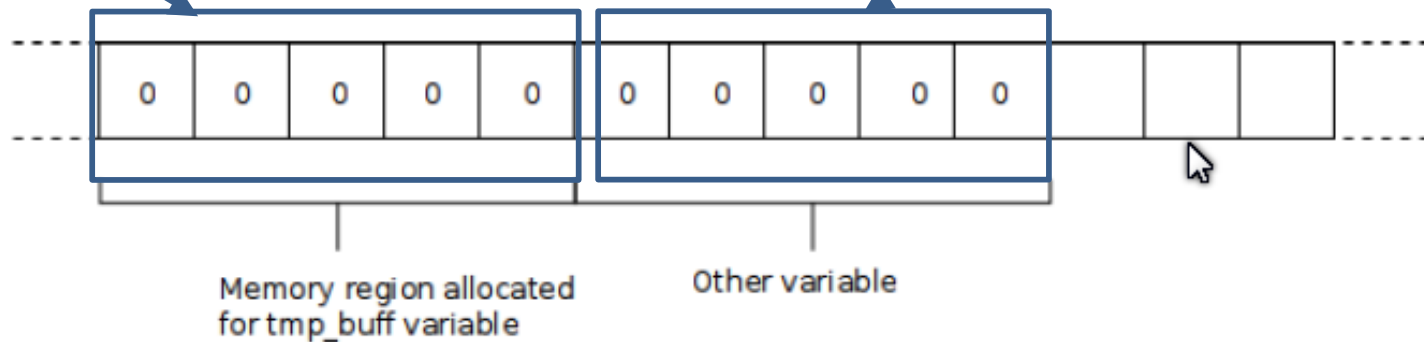
Buffer Overflows: Basics

```
char src[]="ABCDEFGH";
```

```
char tmp_buff[5];
```

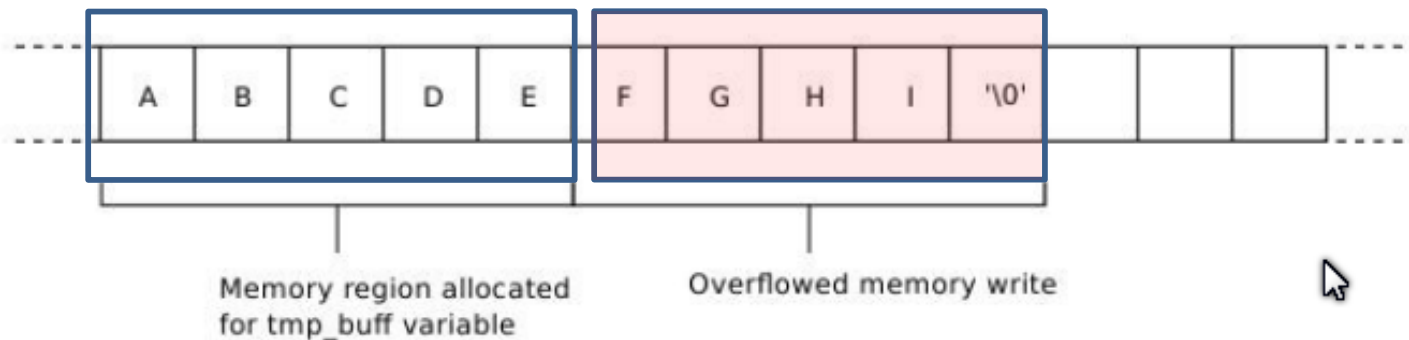
```
int password_checked;
```

```
strcpy(tmp_buff, src);
```



Buffer Overflows: After strcpy()

```
char src[]="ABCDEFGH";  
char tmp_buff[5];  
int password_checked;  
  
strcpy(tmp_buff, src);
```



Buffer Overflows: usual suspects

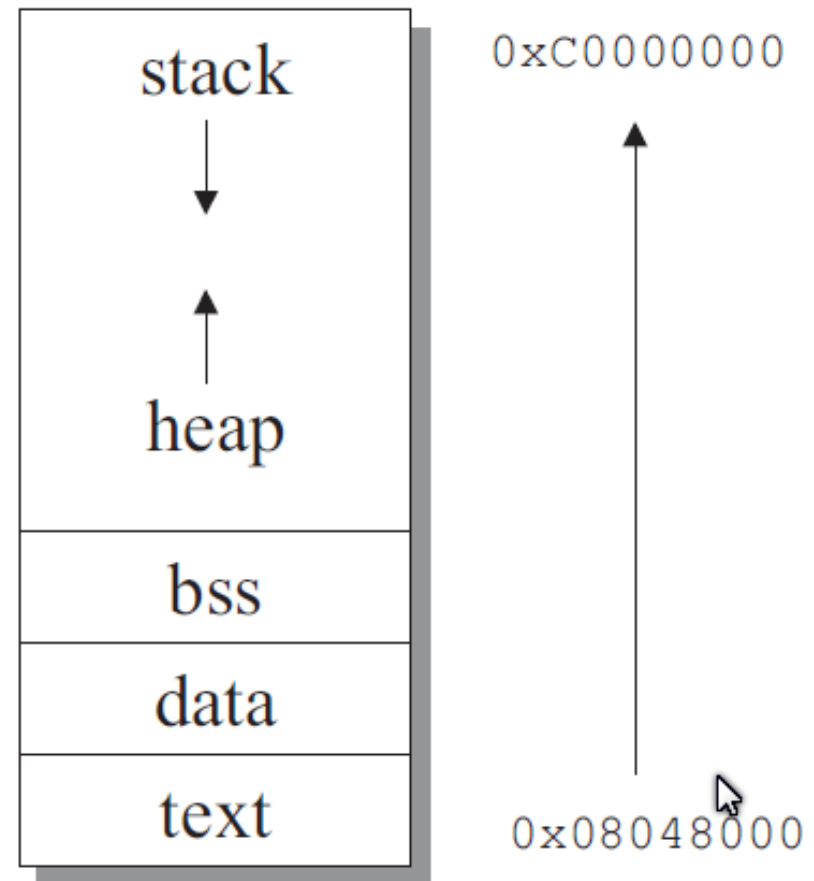
- String manipulation functions that don't properly check string length
 - `gets()`, `strcpy()`, ...
- Copy with incorrect parameters
 - `memcpy()` ...
- Incorrect computation of required memory length
 - zero-sized `malloc()`s ...

Preventing Buffer Overflows

- Safe usage of copy primitives using libc methods
 - For instance, limit amount of data to copy
 - `strncpy(tmp_buff, src, sizeof(tmp_buf))`
- Compiler extension if array length can be computed
- Annotations (e.g. Deputy, CCured...)

More about Memory Layout

- Text
 - Also called code segment
- Data
 - Global initialized data
- BSS
 - Global uninitialized data
- Stack
 - Local variables
 - Also used to store function environments and parameters during calls (stack frames)
 - LIFO
 - Multi-threading : multiple stacks
- Heap
 - Dynamically allocated variables
 - Reserved through `calloc()` and `malloc()`



Program Memory Stack

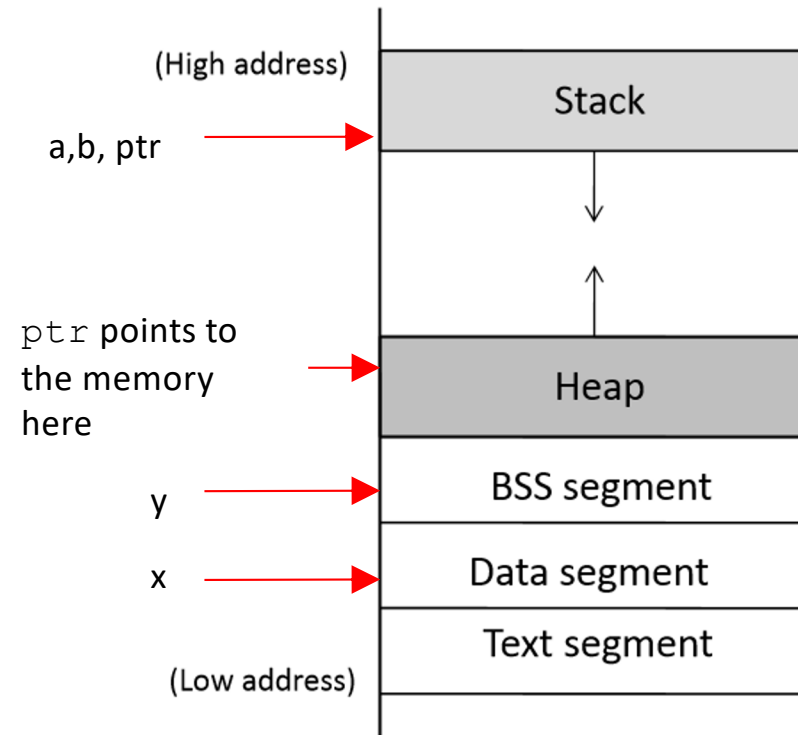
```
int x = 100;
int main()
{
    // data stored on stack
    int a=2;
    float b=2.5;
    static int y;

    // allocate memory on heap
    int *ptr = (int *) malloc(2*sizeof(int));

    // values 5 and 6 stored on heap
    ptr[0]=5;
    ptr[1]=6;


    // deallocate memory on heap
    free(ptr);

    return 1;
}
```




Function Call Stack

```
void f(int a, int b)
{
    int x;
}
void main()
{
    f(1,2);
    printf("hello world");
}
```



Stack Layout for Function Call Chain



- Stack Pointer (`%esp`, `%rsp` ...)
- Base Pointer / Frame pointer (`%ebp`, `%rbp`)
- Instruction Pointer (`%eip`, `%rip`)

Order of the function arguments in stack / offsets (x86/x64)

```
void func(int a, int b)
{
    int x, y;

    x = a + b;
    y = a - b;
}
```

```
movl    12(%ebp), %eax    ; b is stored in %ebp + 12
movl    8(%ebp), %edx     ; a is stored in %ebp + 8
addl    %edx, %eax
movl    %eax, -8(%ebp)    ; x is stored in %ebp - 8
```

Stack Overflow

- Stack Overflows can be caused by:
 - Recursive calls
 - Reentrant interrupts
 - Allocations on the stack
 - Large
 - Controlled by the attacker
 - `Alloca()`, `char array[function_parameter]`

Stack Overflow

- This should never happen ?
- Detection : guard page (Unix)
 - Problem when allocation is too large
 - the stack pointer can “jump over” the guard page
- Prevention: Not so easy
 - Abstract interpretation [*Regehr05*]
 - Works as long as control flow can be determined statically
 - Forbid / bound recursion
 - Forbid / limit allocation on the stack
- Problem with micro-controllers
 - No MMU (Memory Management Unit)

Stack Buffer Overflow / Overrun

- A special case of buffer overflow
 - Exploitation of existing vulnerability in order to modify control flow
- The overflowed buffer is allocated on the stack, typically corrupts return address
- First wide-scale exploitation: Morris' worm (1989) in Unix's fingerd

Vulnerable Program

```
int main(int argc, char **argv)
{
    char str[400];
    FILE *badfile;

    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 300, badfile);
    foo(str);
    printf("Returned Properly\n");
    return 1;
}
```

- Reading 300 bytes of data from badfile.
- Storing the file contents into a str variable of size 400 bytes.
- Calling foo function with str as an argument.

Note : Badfile is created by the user and hence the contents are in control of the user.


Vulnerable Program

```
/* stack.c */
/* This program has a buffer overflow vulnerability. */
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int foo(char *str)
{
    char buffer[100];

    /* The following statement has a buffer overflow problem */
    strcpy(buffer, str);

    return 1;
}
```

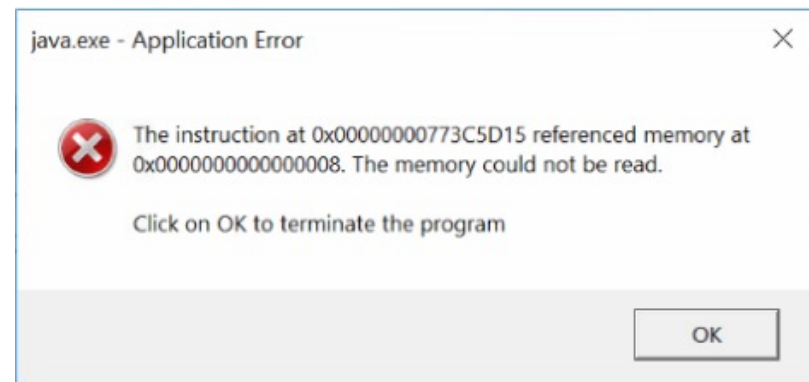


Consequences of Buffer Overflow

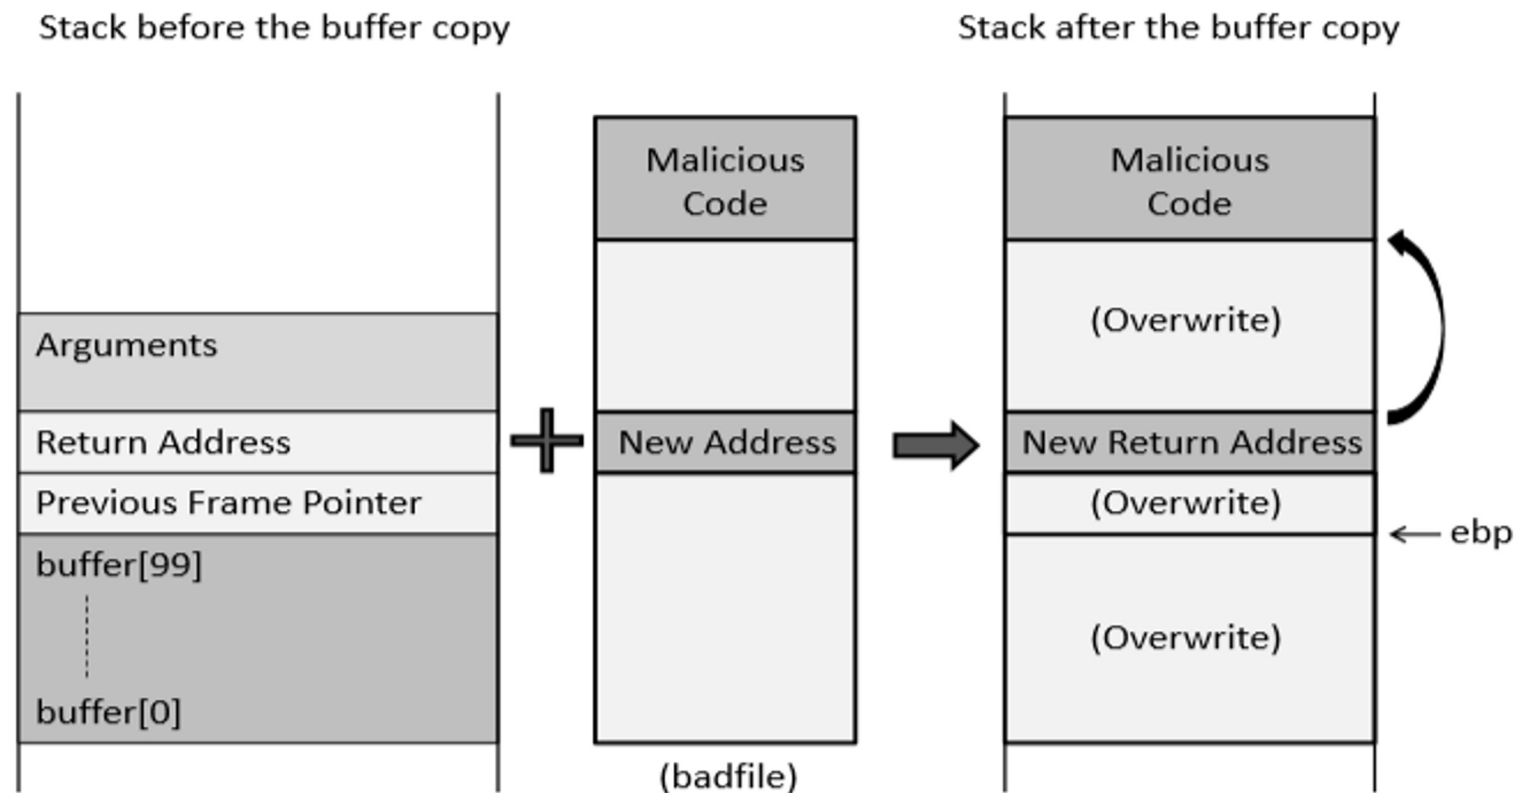
Overwriting return address with some random address can point to :

- Invalid instruction
- Non-existing address
- Access violation

• **Attacker's code** —————> **Malicious code to gain access**

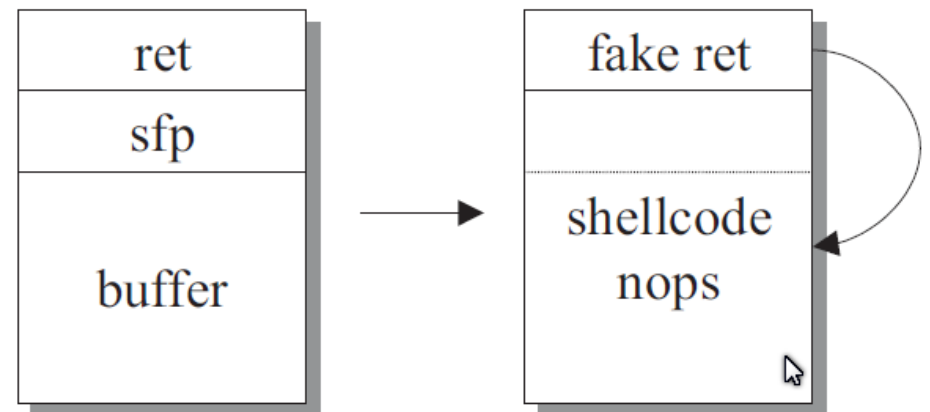
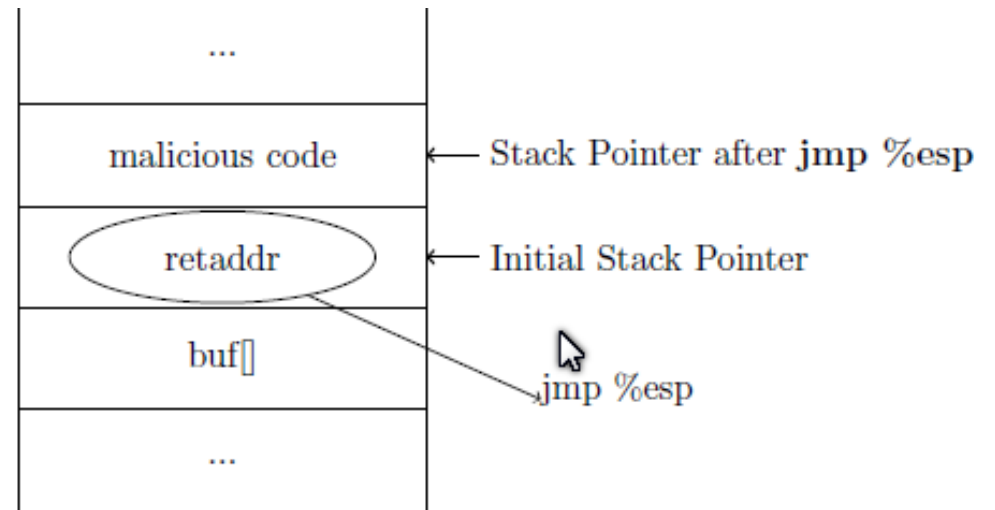


How to Run Malicious Code



More on Modifying return address

- Not always easy to know the exact address
- Trampolines :
 - Return address pointing to program fragment containing a `jmp %esp`
- NOP sled
 - A long sequence of NOP
 - Followed by the shellcode
 - Jumping anywhere in there leads to the shellcode

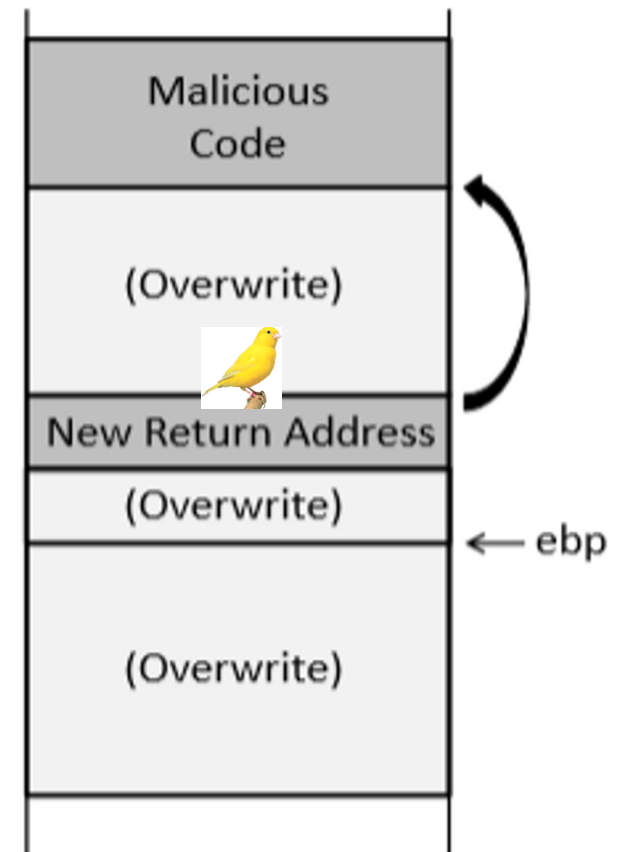


Exploiting Stack Buffer Overflows

- Corrupting the Control Flow
 - Return address in a stack frame
 - Function pointer
 - Exception handler
 - Global Offset Table (GOT)
 - C++ objects
- Corrupting the Program State (other stack frames)
 - Variable that stores authentication status
 - Modify stored password
 - Pointer value

Protections: Canaries

- Objective : Detect unexpected modifications of values on the stack (e.g., return address)
- Inserting a known value on the stack
 - the “canary”
- Compiler extension - inserts code that:
 - Adds a random value after the return address
 - Checks canary value before using return address
- Limitations
 - Canaries can be guessed or obtained through memory leaks
 - Canary copy needs to be stored in secure storage (hopefully not corrupted, or at least makes things harder ...)



Protections: Non Executable Memories (Writable XOR Executable / NX)

- Executable and writable memories regions
 - Allows to directly write instructions (in a buffer) and execute them, e.g., from the stack.
 - Makes stack-based buffer overflows easy to exploit
- W XOR X (A.k.a NX / XN / DEP / PAX)
 - OS splits memory regions – no region is both writable and executable
- Most systems now support a form of W XOR X:
 - Was complex on x86 (segmentation)
 - HW support for it since ~2004
- Many names for the same feature :
 - AMD : Enhanced Virus Protection,
 - ARM : XN eXecute Never
- Defeated by Return-to-libc attacks / Return Oriented programming

Protections: ASLR (Address Space Layout Randomization)

- Randomizes base addresses of memory segments
 - Addresses change at every execution
- Randomize start or base address of:
 - Program code
 - Libraries code
 - Heap/Stack/Data regions
- Objectives:
 - Difficult to guess the stack address in the memory.
 - Difficult to guess %ebp address and address of the malicious code
- Many programs have problems with that
 - Sometimes rewriting part of them is necessary, notably when they contain specially crafted assembly code
 - Code needs to be position independent for program randomization
- Limitations
 - Memory leaks used to learn memory layout
 - Address space / system limitations (e.g., page boundaries) may allow brute-force probes

Address Space Layout Randomization

```
#include <stdio.h>
#include <stdlib.h>

void main()
{
    char x[12];
    char *y = malloc(sizeof(char)*12);

    printf("Address of buffer x (on stack): 0x%x\n", x);
    printf("Address of buffer y (on heap) : 0x%x\n", y);
}
```

Address Space Layout Randomization : Working

```
$ sudo sysctl -w kernel.randomize_va_space=0
kernel.randomize_va_space = 0
$ a.out
Address of buffer x (on stack): 0xbffff370
Address of buffer y (on heap) : 0x804b008
$ a.out
Address of buffer x (on stack): 0xbffff370
Address of buffer y (on heap) : 0x804b008
```

1

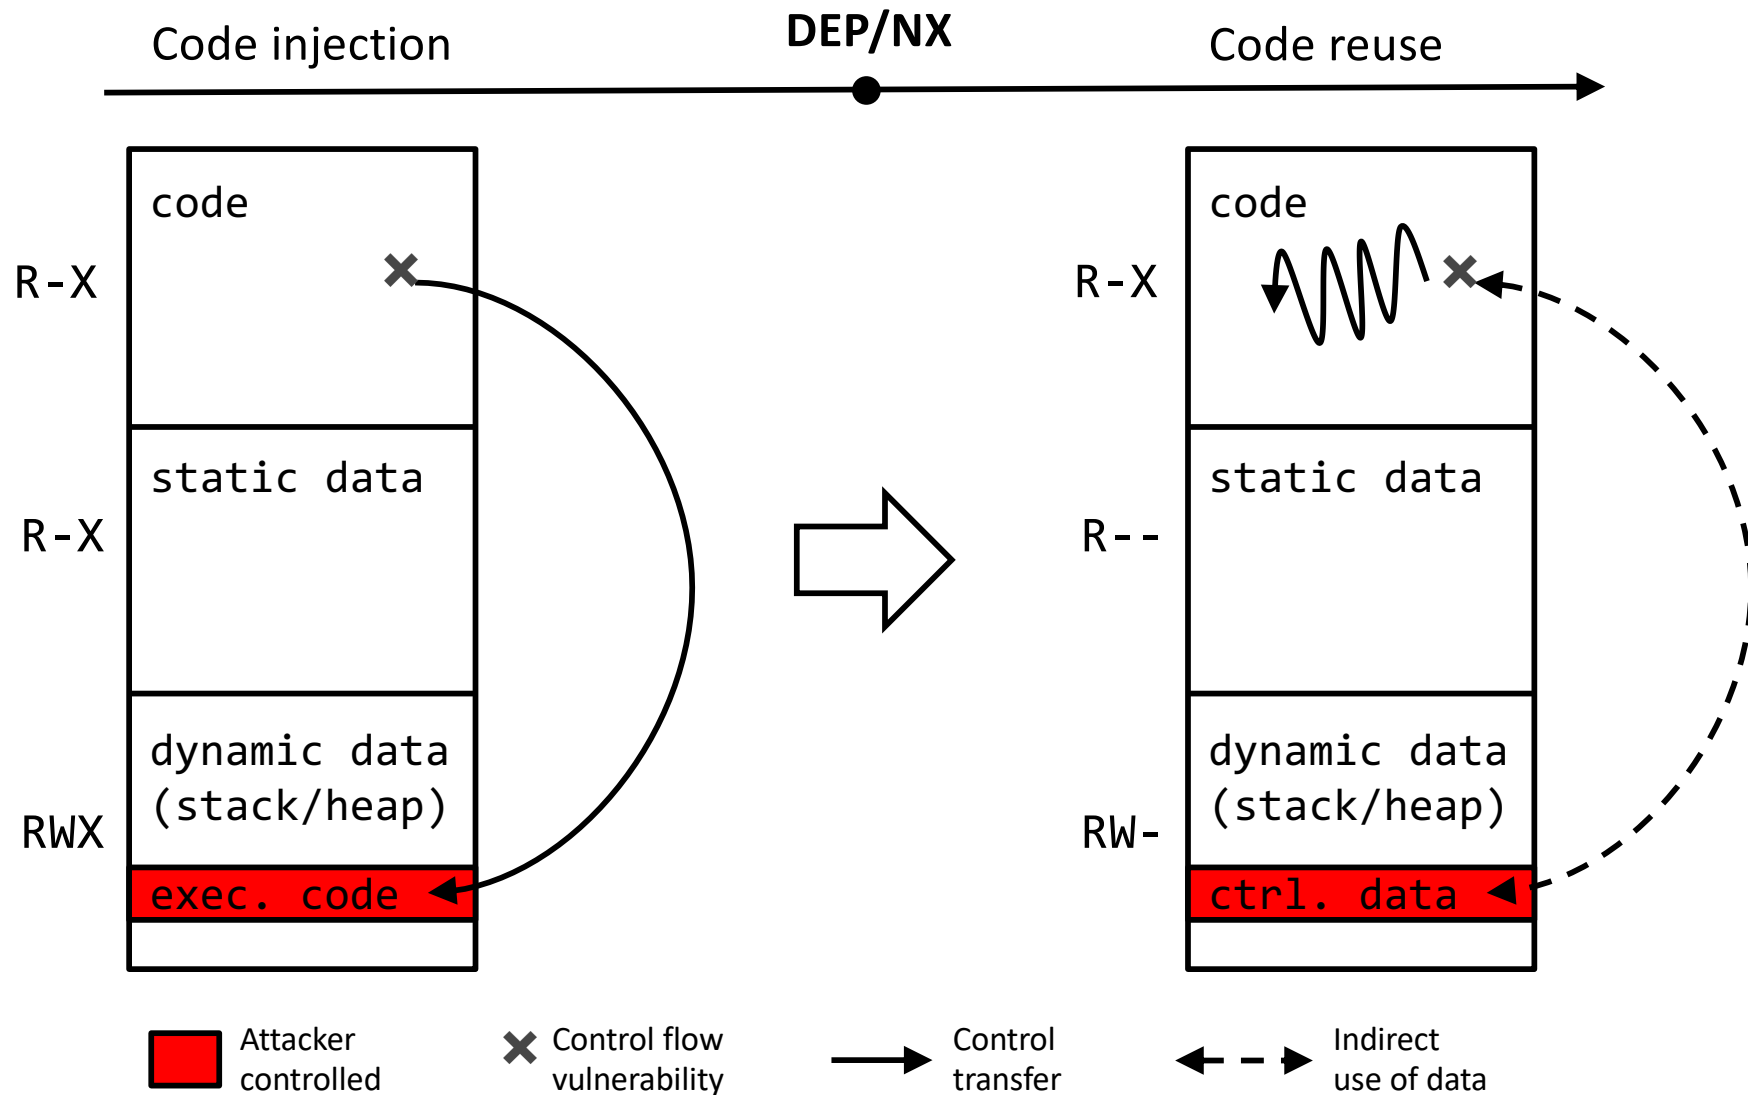
```
$ sudo sysctl -w kernel.randomize_va_space=1
kernel.randomize_va_space = 1
$ a.out
Address of buffer x (on stack): 0xbf9deb10
Address of buffer y (on heap) : 0x804b008
$ a.out
Address of buffer x (on stack): 0xbf8c49d0
Address of buffer y (on heap) : 0x804b008
```

2

```
$ sudo sysctl -w kernel.randomize_va_space=2
kernel.randomize_va_space = 2
$ a.out
Address of buffer x (on stack): 0xbf9c76f0
Address of buffer y (on heap) : 0x87e6008
$ a.out
Address of buffer x (on stack): 0xbfe69700
Address of buffer y (on heap) : 0xa020008
```

3

Evolution of machine code level attacks



Return-to-libc

- NX (W XOR X) makes it impossible to inject one's code and execute it.
 - No memory regions that are writeable and executeable
- Idea : Reuse existing code
 - “Fortunately” libc loaded at a constant address
 - Divert control flow of exploited program into libc code
 - “Load” parameters on the stack
 - No code injection required: Jump to a known address
 - `exec()`, `system()`, `printf()`
- For example:
 - `Exec(“/bin/sh”)`

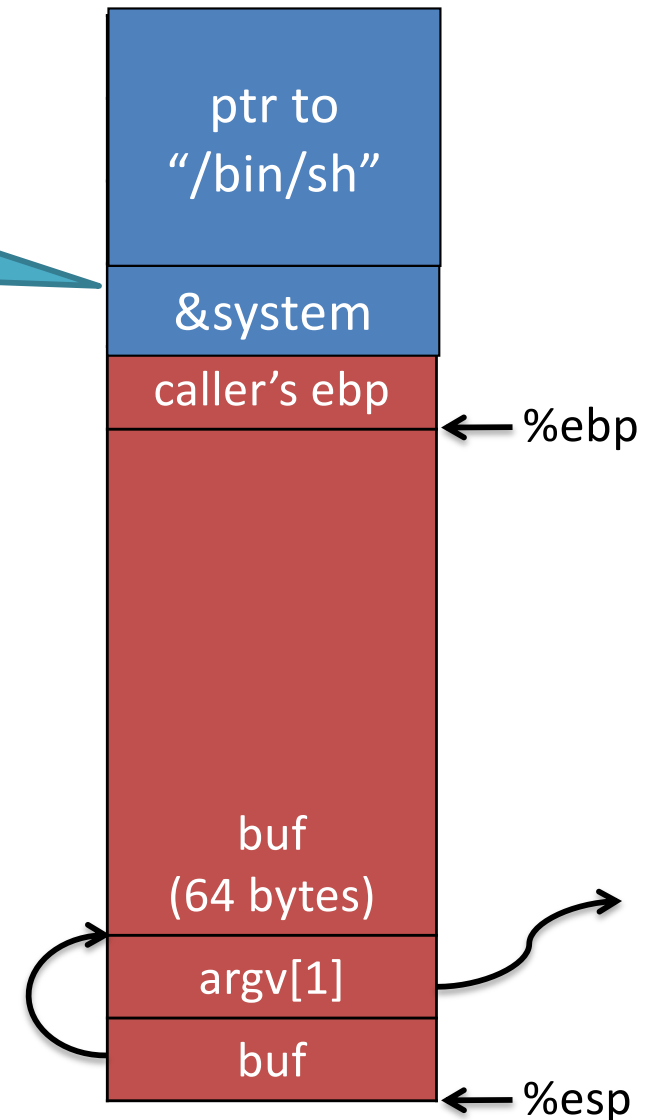
Return-to-libc Attack

ret transfers control to system,
which finds arguments on stack

Overwrite return address with
address of libc function

- setup fake return address and
argument(s)
- ret will “call” libc function

No injected code!



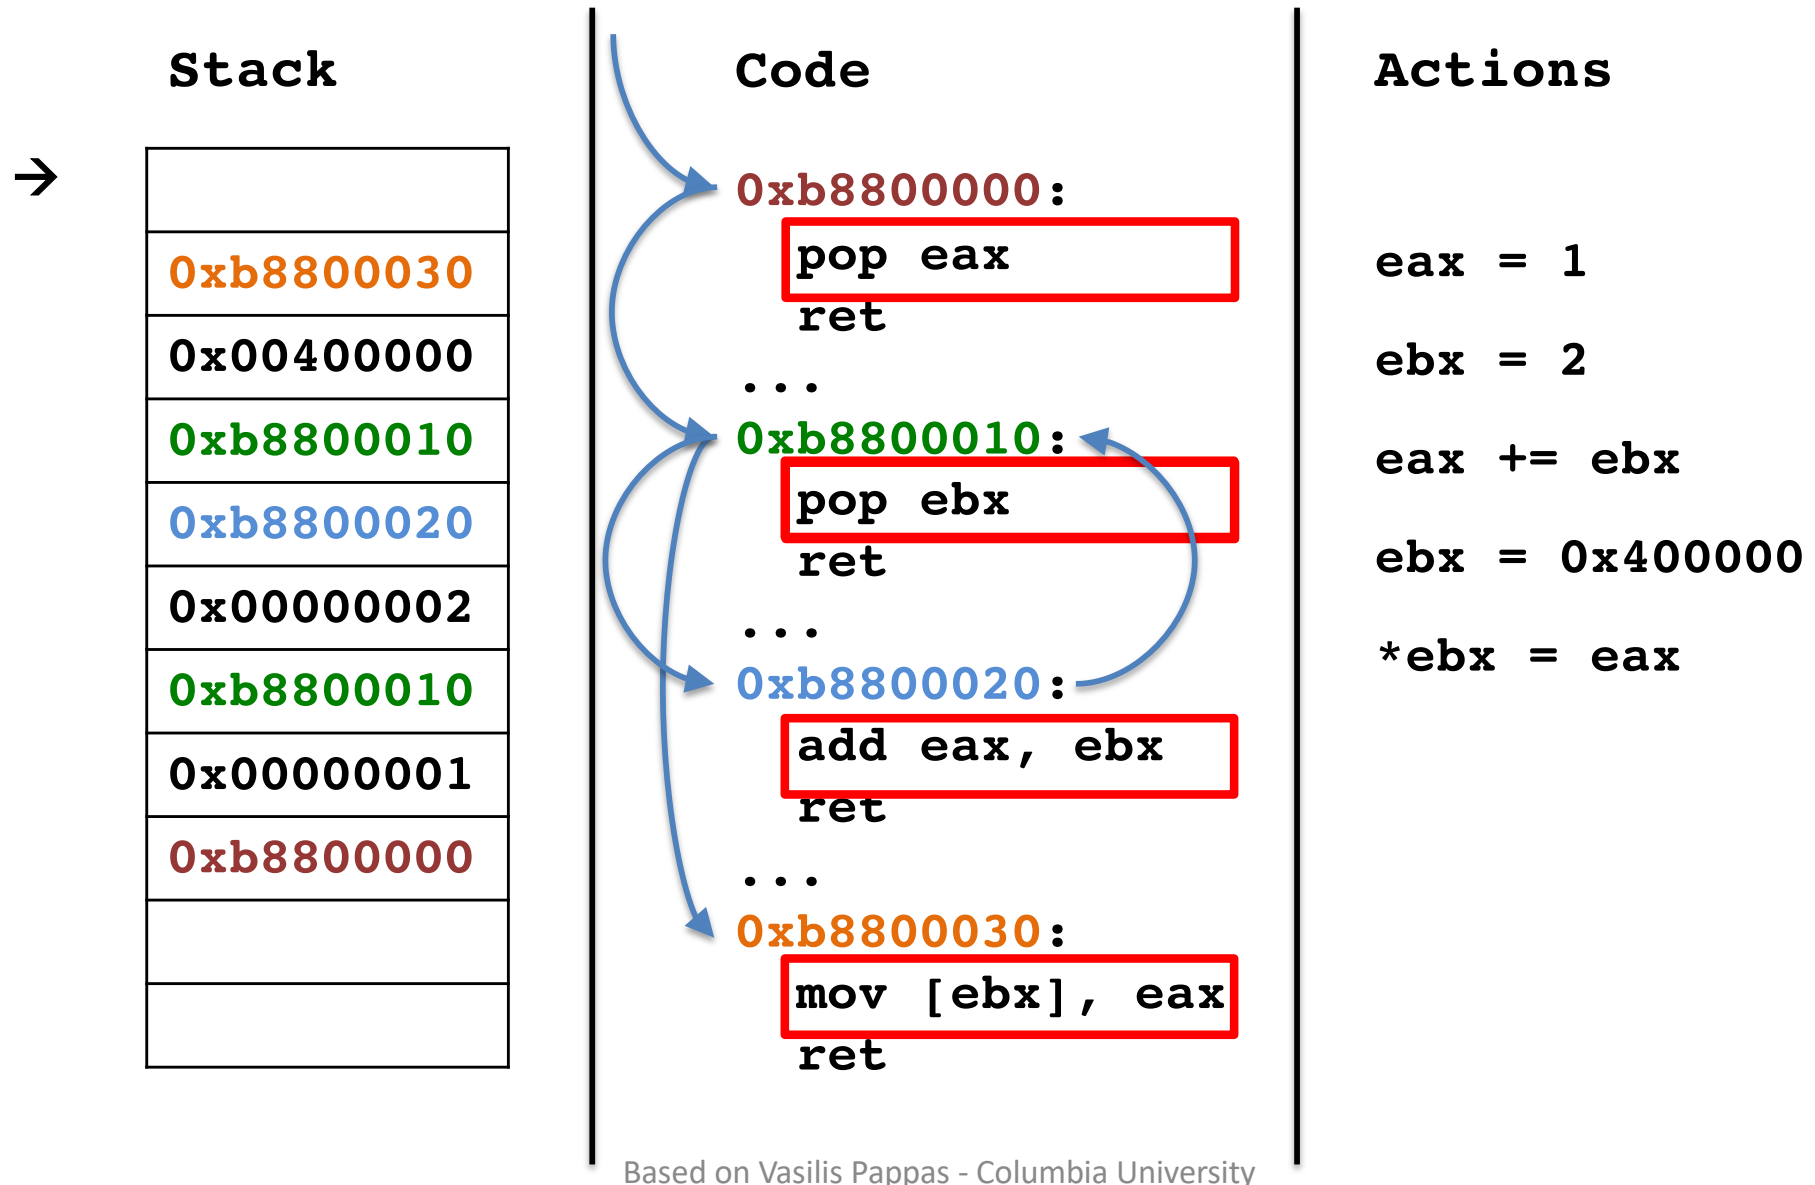
Return-Oriented Programming (ROP)

- return-into-libc seems limited and easy to defeat
 - Attacker cannot execute arbitrary code
 - Attacker relies on contents of libc
- This perception is false: Return-Oriented Programming & Jump-Oriented Programming
 - A special case of return-into-libc
 - Arbitrary attacker computation and behavior (given any sufficiently large codebase to draw on)

ROP: Approach

- Most directly inspired by *Borrowed code chunks* [Krahmer 2005]
 - Find short sequences of instructions that allow to perform some given operations
 - Termed Gadgets
 - “Chain” them together using “ret”
- JOP attack = use *jmp* instead of *ret*

Return-Oriented Programming



ROP: Approach

- A Turing complete set of gadgets allows to perform arbitrary computation
 - Exploits are not straight-line limited
 - Showed to work on most architectures
 - Equivalent to having a virtual machine/interpreter
- Calls no functions at all
 - can't be defeated by removing functions like `system()`
 - Must know the memory map (no ASLR)
 - Need to find interesting gadgets and to chain them in a given order
- Specific compilers (e.g. ROPC)
 - Automation techniques to find those sequences of code
 - Satisfiability Modulo Theories (SMT) Solvers

ROP: consequences & protection


- Malicious code detection cannot be limited to executable memory regions
 - Return oriented rootkits / malicious code...
 - Even non executable memories need to be verified
- ROP defeated by ASLR
 - chaining returns needs to know addresses in advance
- Blind ROP
 - It is possible to learn where the gadgets are, brute force and monitor side effects
 - Stack learning overwrite a byte at a time and brute force it.

Heap Buffer Overflows

- The heap is the pool of memory used for dynamic allocations at runtime
 - malloc() grabs memory on the heap
 - free() releases memory on the heap

- Blocks of data are stored in a doubly linked list

```
typedef struct __HeapHdr__ {  
    struct __HeapHdr__ *next;  
    struct __HeapHdr__ *prev;  
    unsigned int size;  
    unsigned int used;  
    // Usable data area starts here  
} HeapHdr_t;
```



Heap Buffer Overflow Vulnerability

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>


int main(int argc, char **argv) {
    char *p, *q;

    p = malloc(1024);
    q = malloc(1024);
    if (argc >= 2)
        strcpy(p, argv[1]);
    free(q);
    free(p);
    return 0;
}
```

```
% ./heapbug `perl -e 'print "A"x5000'`
Segmentation fault
```

Heap Buffer Overflows

- next/prev pointers are stored after the data
 - Overflow: overwrite the prev/next pointers (headers)
- Freeing a chunk = update double linked list
 - This allows arbitrary value write at arbitrary address (red = attacker controlled), e.g. function pointer
 - FD = hdr -> **next**
 - BK = hdr -> **prev**
 - **FD**->prev = **BK**
 - **BK**->next = **FD**
 - Cf. <https://www.win.tue.nl/~aeb/linux/hh/hh-11.html>
- Detection is simple:
 - Test if (hdr->prev->next == hdr) otherwise an attack is underway!
 - canaries



Heap Overflow Exploitation

- Direct attacks: modify function pointer
 - Simple overflow to the pointer location
- Often indirect attacks on the stack return address
 - Fill headers with the address of the return address on the stack
 - The next malloc/free operation will modify the return address at will
- Heap spraying:
 - Exploits contiguous chunk placement (e.g., browser, PDF, Flash)
 - Fill up an entire chunk with NOP sled + shellcode and spray it repeatedly into the heap
- Can be very complex
 - Need to predict heap layout, control program state
 - Otherwise lead program in a state where it is exploitable